

**An Analytical Study of AI-Driven Process Engineering for Automation and
Optimization of IT Operations (AIOps)**

Shreevikas Bangalore Jagadish (A20572109)

Department of Information Technology and Management, Illinois Institute of Technology

ITMM 572: Process Engineering for IT Managers

Professor Nazneen Hashmi

March 26, 2026

An Analytical Study of AI-Driven Process Engineering for Automation and Optimization of IT Operations (AIOps)

Introduction

Modern IT operations no longer revolve around a few stable servers and a manageable stream of manually triaged tickets. Today's environments combine public cloud, private infrastructure, SaaS platforms, containers, microservices, APIs, CI/CD pipelines, observability tools, and security controls. As these ecosystems have become more dynamic and interdependent, the volume of telemetry and the rate of operational change have exceeded what human operators can interpret reliably through manual inspection alone. AIOps emerged as a response to this overload by applying analytics, machine learning, and increasingly large language model based reasoning to operational data so organizations can detect issues earlier, diagnose them faster, and respond more consistently (Cheng et al., 2023; Notaro et al., 2021).

Although AIOps is often marketed as a technology category, its importance is better understood as a management and process question. The real value does not come merely from a smarter anomaly detector or a more sophisticated chatbot. It comes from redesigning the workflow through which an organization observes a system, interprets signals, decides what matters, routes work, executes actions, and learns from outcomes. That is why process engineering offers such a strong analytical lens. Process engineering examines the sequence of activities, the ownership of decisions, the controls that shape action, the information that moves between roles, and the metrics that reveal performance. Under that lens, AIOps is not simply automation; it is AI-driven process redesign for IT operations (Diaz-de-Arcaya et al., 2024).

Academic and practitioner literature supports this broader interpretation. Surveys of AIOps routinely organize the field around lifecycle tasks such as failure detection, root cause

analysis, prioritization, remediation, and learning. Research on microservice reliability also shows that operational performance depends on combining logs, metrics, traces, topology, and change data across distributed systems. More recent studies on LLM-based AIOps go even further by showing how runbooks, tickets, knowledge bases, and incident narratives can be integrated into operational reasoning loops. Across these streams, the common idea is that AIOps turns fragmented evidence into more coordinated operational decisions (Soldani & Brogi, 2023; Zhang et al., 2025).

This paper argues that AIOps is best understood as AI-driven process engineering for the automation and optimization of IT operations. Its strategic value lies in redesigning end-to-end operational flow rather than automating isolated tasks. When implemented well, AIOps can reduce alert noise, shorten incident cycles, remove operational toil, improve service reliability, and strengthen organizational learning. When implemented poorly, it can create opaque automation, weaken accountability, and amplify poor-quality data or weak process design. Therefore, the central thesis of this paper is that AIOps succeeds not because of model novelty alone, but because of strong workflow design, measurement discipline, governance, and human-centered operational integration.

AIOps Through a Process Engineering Lens

Process engineering focuses on how work moves through a system: what inputs trigger activity, what decisions occur, where delays and errors arise, which controls constrain action, and how outcomes are measured. When that lens is applied to IT operations, several chronic inefficiencies become visible. Alert streams are duplicated across tools, engineers spend excessive time searching for context, escalation logic differs by team, and lessons learned from prior incidents are often stored weakly or inconsistently. In many environments, response quality

depends too heavily on individual memory instead of a reliable operating model. These conditions make IT operations an especially strong candidate for process redesign supported by AI (Diaz-de-Arcaya et al., 2024; Google Cloud, 2020).

AIOps addresses those inefficiencies by changing the order and ownership of operational work. Instead of forwarding every alert directly to a human queue, an AIOps workflow can suppress duplicates, correlate symptoms, enrich the issue with topology and change data, estimate likely impact, and route the issue to the appropriate owner or automation path. Instead of requiring engineers to search across dashboards, tickets, and logs manually, the system can construct a contextual incident package with supporting evidence. Instead of handling recurring incidents as isolated events, the workflow can capture patterns and feed them back into runbooks, policies, and future automation. This is process engineering because the primary gain is in workflow quality, not merely in algorithmic sophistication (Notaro et al., 2021).

A useful distinction is the difference between task automation and process optimization. Task automation handles a bounded activity such as classifying a ticket, summarizing an incident, or ranking alerts. Process optimization improves the entire workflow from detection through remediation and review. The latter is strategically more important because isolated automation may save minutes, while workflow redesign can systematically reduce queue time, improve consistency, lower operational error variance, and create repeatable decision pathways. Organizations often overestimate the value of narrow automation because it is visible and easy to demo, while underestimating the broader value of redesigning the operational system itself.

Seen in this way, AIOps belongs with service management, site reliability engineering, operational excellence, and continuous improvement rather than outside them. It is not a substitute for process discipline. Instead, it is a means of embedding intelligence into

observability, triage, coordination, and feedback loops so that operational discipline can scale under conditions of complexity and speed.

Literature Review and Conceptual Foundations

The literature on AIOps has developed from early surveys of failure-management methods into a broader body of work spanning observability analytics, microservice reliability, MLOps integration, and LLM-based operational assistance. Notaro et al. reviewed one hundred AIOps solutions and categorized them into tasks such as failure detection, localization, diagnosis, and remediation. Their work remains foundational because it makes clear that AIOps is not one problem but a chain of interdependent tasks, each with different assumptions, data needs, and evaluation criteria (Notaro et al., 2021).

Cheng et al. expand the discussion by framing AIOps as a cloud-scale capability shaped by the complexity of modern digital platforms. Their review emphasizes that operational intelligence depends on combining heterogeneous data sources and that high rates of system change make static rule systems increasingly insufficient. Soldani and Brogi make a similar point in the microservice context, showing that anomaly detection and root cause analysis are challenging because failures propagate across distributed services and because different telemetry sources tell only partial parts of the story (Cheng et al., 2023; Soldani & Brogi, 2023).

Diaz-de-Arcaya et al. add a valuable managerial dimension by examining the overlap between MLOps and AIOps. Their work argues that successful adoption depends not only on modeling techniques but also on roadmaps, lifecycle control, organizational readiness, and integration between data, models, and operations. This is especially relevant for process engineering because it suggests that AIOps maturity is an organizational capability question as

much as a technical one. The field is therefore moving away from purely model-centric thinking toward lifecycle and operating-model thinking (Diaz-de-Arcaya et al., 2024).

The newest literature focuses on LLMs and agentic systems. Zhang et al. show that large language models extend AIOps beyond narrow prediction tasks toward richer reasoning over incident reports, runbooks, dashboards, chat discussions, and knowledge repositories. This opens new possibilities for explanation, decision support, knowledge retrieval, and orchestrated action. At the same time, it introduces new concerns around hallucination, grounding, overreliance, and safe autonomy. Thus, the literature increasingly converges on a balanced conclusion: AIOps is becoming more capable, but its value depends on guardrails, evaluation, and governance rather than on AI capability alone (Zhang et al., 2025; NIST, 2024).

Process Architecture of AIOps

AIOps can be analyzed as a layered process architecture with four major stages: observability and data acquisition, intelligence generation, decision orchestration, and feedback learning. This layered view is useful because many organizations overfocus on the intelligence stage while underinvesting in data quality, governance, or learning design. In practice, the full control loop determines value (Cheng et al., 2023; Diaz-de-Arcaya et al., 2024).

The first stage is observability and data acquisition. AIOps relies on logs, metrics, traces, events, topology maps, configuration records, change histories, service-desk data, and sometimes business-impact signals. The challenge is not only the volume of data but its fragmentation. Different tools often encode time, severity, ownership, and system identity differently. If telemetry cannot be normalized and linked across systems, AI models inherit ambiguity. Microsoft’s health-modeling and Azure Monitor guidance both suggest that meaningful AIOps

begins with a structured operational baseline rather than with AI layered on top of disordered monitoring data (Microsoft, 2024; Microsoft, 2025b).

The second stage is intelligence generation. This includes anomaly detection, event correlation, root cause inference, severity estimation, summarization, and recommendation generation using statistical methods, machine learning, graph reasoning, and language models. At this stage, the design challenge is not only accuracy but operational usefulness. A model that identifies anomalies without reducing investigation effort may still be weak in practice. Likewise, a diagnosis model that performs well offline but breaks under changing topology or incomplete context can generate false confidence. Research repeatedly shows that AIOps models should be judged by their downstream contribution to operational flow and not only by isolated benchmark scores (Notaro et al., 2021; Soldani & Brogi, 2023).

The third stage is decision orchestration. Once an issue is detected, the system must determine what happens next: who is notified, how severity is assigned, whether remediation is recommended or executed, what approvals are required, and how action is documented. This stage converts analysis into operational impact. Recent Microsoft work on incident optimization and Azure SRE Agent shows that enterprise AIOps is moving toward integrated orchestration in which AI not only interprets incidents but also coordinates incident workflows and, under defined constraints, executes low-risk actions (Microsoft, 2025a; Microsoft, 2025c).

The fourth stage is feedback and continuous learning. Every incident produces evidence about model quality, workflow design, runbook usefulness, and the appropriateness of automation boundaries. Mature AIOps programs feed those outcomes back into thresholds, prompts, playbooks, routing logic, and policy design through post-incident review and periodic

measurement. Without this learning stage, the automation layer stagnates and the same problems recur under new forms (Microsoft, 2026a; NIST, 2023).

Core Application Areas

The most established AIOps application is incident detection and triage. Modern platforms generate enormous volumes of alerts, many of which are low-value, duplicate, or symptoms of the same underlying problem. AI helps by clustering similar events, suppressing noise, ranking urgency, and enriching issues with context. The major process benefit is fewer handoffs and faster prioritization. In overloaded on-call environments, this can materially improve both speed and operator focus (Cheng et al., 2023; Notaro et al., 2021).

A second major application is root cause analysis. In distributed systems, visible symptoms often appear far from the originating fault. Metrics may show downstream degradation, traces may reveal latency propagation, and logs may show only fragments of the true cause. Survey work on microservices consistently shows that diagnosis is strongest when multiple signals are combined and interpreted in the context of service relationships. AIOps is therefore especially valuable when it can move beyond isolated anomaly flags and provide a defensible explanation of likely failure propagation (Soldani & Brogi, 2023).

A third application is automated remediation. Once organizations have sufficient trust in detection and diagnosis quality, they can automate low-risk actions such as restarting services, clearing queues, triggering scale adjustments, rolling back known-bad changes, or opening structured incident workflows with relevant evidence attached. This is where AIOps shifts from assistive intelligence to operational execution. However, this is also where governance becomes more demanding because automation must be matched to risk level, reversibility, auditability, and human override mechanisms (Microsoft, 2025c; Microsoft, 2026b).

A fourth application is performance, capacity, and cost optimization. AIOps can identify underutilized or overstrained resources, forecast saturation risk, recommend right-sizing, and support planning decisions that balance reliability and financial efficiency. This matters especially in cloud environments where performance and cost are tightly coupled. AIOps therefore extends beyond incident response into ongoing operational optimization (Cheng et al., 2023; Microsoft, 2025b).

A fifth application is knowledge augmentation. LLM-enabled systems can summarize incidents, retrieve relevant runbooks, answer diagnostic questions, and make operational knowledge more accessible to less experienced engineers. In process terms, this turns documentation from passive storage into active decision support. The value is not just speed but consistency, because operational quality becomes less dependent on who happens to be on shift (Zhang et al., 2025).

Benefits for Automation and Optimization of IT Operations

A well-engineered AIOps capability produces benefits across speed, quality, resilience, and labor efficiency. The most immediate benefit is faster detection and response. If issues are identified earlier and routed with better context, teams spend less time on uncertainty and more time on action. In many incidents, the first delay is not the repair itself but the effort required to determine what is happening, who should respond, and how severe the situation is. AIOps directly targets that delay (Microsoft, 2025a; Notaro et al., 2021).

A second major benefit is the reduction of operational toil. Google's SRE guidance defines toil as work that is manual, repetitive, automatable, tactical, and low in enduring value. Many operations tasks fit that definition: clearing repetitive alerts, opening routine tickets, running the same checks, or repeating standard diagnostics. AIOps can remove a substantial

portion of that burden, allowing engineering time to shift from repetitive maintenance toward structural reliability improvement (Google Cloud, 2020).

A third benefit is improved organizational learning. When incidents are processed through a workflow that captures context, hypotheses, actions, approvals, and outcomes, the organization gradually builds a reusable knowledge asset. This strengthens postmortems, improves training of new staff, and reduces dependence on tribal memory. Over time, each incident can improve future readiness rather than disappearing as an isolated event (Diaz-de-Arcaya et al., 2024; Microsoft, 2026a).

A fourth benefit is better business alignment. IT operations is often measured using purely technical indicators, yet operational issues also affect customer experience, revenue exposure, compliance posture, and reputation. AIOps can help connect incident handling to service-level objectives, mitigation value, backlog reduction, and avoided downtime. This is essential if organizations want to evaluate AIOps as a business capability and not merely a technical experiment (Microsoft, 2026a; NIST, 2023).

Implementation Framework and Maturity Model

Organizations should not begin their AIOps journey with the most ambitious autonomous use cases. A safer and more effective path is progressive maturity. The first phase is observability readiness. Teams need adequate instrumentation, consistent identities across systems, usable topology, clean timestamping, and governance over critical telemetry sources. Without these foundations, later AI stages inherit data confusion and produce unreliable outputs (Cheng et al., 2023).

The second phase is process mapping. Teams should document the current-state workflow for detection, triage, diagnosis, approval, remediation, and post-incident review. The

purpose is to identify where delay, ambiguity, duplication, and inconsistent decision making occur. This step matters because AIOps should improve the process, not simply automate an already broken workflow (Diaz-de-Arcaya et al., 2024).

The third phase is assistive AIOps. Here, AI enriches incidents, retrieves relevant knowledge, ranks possible causes, and recommends next steps, but humans remain the primary decision makers. This phase is useful because it creates measurable operational gains while also generating evidence about where the system is accurate, where it is weak, and which parts of the workflow are most suitable for further automation.

The fourth phase is bounded automation. Organizations identify low-risk, high-frequency actions and define explicit controls for them. A system may, for example, restart a stateless service or open a standard incident workflow when specific conditions are met, while higher-risk actions still require approval. This layered autonomy model is stronger than the idea of full automation because it ties authority to risk level and control maturity (Microsoft, 2025c; NIST, 2023).

The fifth phase is continuous optimization. At this point, AIOps becomes part of the operating model. Metrics are reviewed regularly, post-incident findings update policies and automation, and governance processes monitor drift, exceptions, and value realization. Maturity is reached when AIOps is managed as an evolving operational capability rather than as a one-time project.

Measurement and KPI Design

Measurement is central to AIOps because systems can appear impressive while producing weak operational value. A mature KPI design should therefore include technical, workflow, business, and governance metrics. Technical metrics may include anomaly precision, false-positive rate, route accuracy, recommendation acceptance rate, and stability across changing environments. These are important but not sufficient on their own (Notaro et al., 2021; Soldani & Brogi, 2023).

Workflow metrics reveal whether the process is actually improving. Typical examples include mean time to detect, mean time to acknowledge, mean time to resolve, reduction in duplicate tickets, escalation frequency, autonomous mitigation rate, and percentage of incidents resolved without human paging. These indicators show whether AIOps is shortening response loops and reducing labor friction instead of merely shifting work elsewhere (Google Cloud, 2020; Microsoft, 2026a).

Business metrics connect operational improvement to managerial outcomes. These may include avoided downtime, severity mix, SLA or SLO attainment, customer-impact minutes, backlog reduction, and cloud spend optimization. Many AIOps programs struggle to retain executive sponsorship because they report only model scores rather than operational or financial value. Strong measurement closes that gap.

Governance metrics matter just as much. Organizations should track override rates, rollback frequency, policy violations, audit completeness, drift in model behavior, and the frequency of unsupported or low-confidence outputs in LLM-assisted workflows. These measures help ensure that speed does not come at the expense of control and accountability.

Human Factors and Organizational Design

AIOps is often discussed in technical language, but its effectiveness depends heavily on human factors. Operators must understand what the system is doing, when to trust it, and when to override it. If recommendations are opaque, teams may ignore them. If automation is too aggressive or badly explained, teams may resist adoption or create side processes outside the approved workflow. Trust depends on evidence, explainability, and consistency rather than on vague claims of intelligence (NIST, 2023).

There is also a skill-design question. As AIOps reduces repetitive manual work, the role of operations engineers shifts. Less time may be spent manually correlating alerts, and more time may be spent designing controls, maintaining runbooks, validating models, interpreting complex exceptions, and improving architecture. That is beneficial, but only if organizations support the transition with role clarity and training. Otherwise, AIOps can create uncertainty, resistance, or fears of deskilling (Diaz-de-Arcaya et al., 2024).

Knowledge management is another human factor. In many organizations, valuable operational knowledge is hidden across tickets, chat channels, dashboards, and the memory of senior staff. LLM-enabled AIOps can surface that knowledge more effectively, but only when the underlying sources are curated, current, and governed. Poorly maintained documentation degrades AI assistance just as poor telemetry degrades anomaly detection (Zhang et al., 2025).

Finally, change management matters. AIOps adoption works best when engineers are involved in defining requirements, reviewing boundaries, and evaluating outcomes. Introduced that way, AIOps feels less like a black-box replacement and more like a disciplined redesign of operations.

Governance, Risk, and Responsible AIOps

As AIOps systems gain more autonomy, governance becomes a central concern. The NIST AI Risk Management Framework is highly relevant because it emphasizes governing AI systems, mapping context, measuring risk, and managing the system across its lifecycle. Although the framework is not AIOps-specific, its principles apply directly to operational automation that can affect service availability, security, compliance, and cost. AIOps systems therefore require documented ownership, approved boundaries, monitoring procedures, and response plans for AI-related failures (NIST, 2023).

Generative and agentic systems introduce additional risks. The NIST Generative AI Profile highlights issues such as confabulation, overreliance, unsafe output, and failures in transparency or accountability. In an AIOps context, these translate into unsupported remediation suggestions, misleading summaries, actions based on weak grounding, or misplaced trust in fluent explanations. In operational settings, such failures can have direct business consequences, which means they must be governed actively rather than treated as acceptable experimentation (NIST, 2024).

Practitioner documentation points in the same direction. Azure SRE Agent materials emphasize approvals, auditability, incident-value tracking, and logging of agent actions. This is significant because it shows that serious enterprise AIOps is evolving around concrete operational controls: who can authorize actions, what the system did, why it acted, and what happened afterward. Governance becomes operationally meaningful when it is visible in the workflow itself rather than in separate policy documents (Microsoft, 2026a; Microsoft, 2026b).

Responsible AIOps therefore requires a layered control model: role-based access, environment scoping, approval gates for high-risk actions, mandatory logging of tool use and

decisions, regular review of failures and drift, and clear incident response procedures for the AI system itself. The more authority the system has, the stronger the surrounding controls must be.

Large Language Models and Agentic AIOps

Large language models are changing AIOps because they can work across unstructured and semi-structured operational information in ways earlier systems could not. Traditional AIOps tools were strongest on pattern detection in logs, metrics, or events but weaker at using runbooks, tickets, architectural notes, or conversational context. LLMs make it possible to combine those sources into richer operational reasoning, explanation, and summarization workflows (Zhang et al., 2025).

This creates an important process shift. Earlier AIOps pipelines often stopped at identifying an anomaly or ranking a likely cause. Agentic AIOps can potentially continue the workflow by retrieving relevant documentation, generating a mitigation plan, checking permissions, calling tools, and documenting the result. Microsoft's Azure SRE Agent reflects this movement by positioning the agent as an AI-powered reliability assistant that helps investigate incidents and automate remediation tasks across Azure resources while preserving approval and audit controls (Microsoft, 2025c; Microsoft, 2026b).

However, language models are persuasive generators, not inherently dependable operational reasoners. Their usefulness depends on grounding, action scoping, evaluation, documentation quality, and explicit uncertainty handling. In AIOps, a fluent but unsupported answer can be more dangerous than no answer at all because operators may over-trust it. Retrieval grounding, tool constraints, evidence display, and human review therefore remain essential design choices (NIST, 2024; Zhang et al., 2025).

The most realistic near-term outcome is layered autonomy. LLMs and agents will handle high-volume synthesis, context assembly, and bounded actions, while humans retain responsibility for ambiguous cases, high-risk decisions, policy setting, and exception handling. This division of labor aligns with both process engineering and responsible AI principles.

Critical Limitations and Trade-Offs

Despite its promise, AIOps can fail in predictable ways. One failure mode is poor data quality. If service identities are inconsistent, topology is outdated, timestamps drift, or change records are incomplete, the system may produce confident-looking but operationally weak outputs. Process engineering helps expose this by forcing attention to information flow quality rather than assuming the model alone can compensate for structural data issues (Cheng et al., 2023).

A second failure mode is automation without control. Organizations may automate aggressively to achieve speed metrics, but automation that lacks rollback paths, auditability, policy enforcement, or clear ownership can amplify incidents instead of resolving them. The relevant managerial trade-off is not automation versus no automation, but speed versus controlled reliability. High-performing AIOps programs recognize that not every task should be autonomous and that autonomy should be shaped by reversibility, blast radius, and evidence quality (NIST, 2023).

A third failure mode is overreliance on LLMs. Agentic systems create new attack surfaces and trust problems, including prompt injection through operational data, unsupported recommendations, or hallucinated remediation steps. Even when malicious actors are absent, model outputs may still be overly confident or insufficiently grounded. That means prompt

security, data sanitization, action constraints, and fallback behavior are operational requirements rather than optional add-ons (NIST, 2024).

A fourth failure mode is fragmented ownership. AIOps usually crosses the boundaries of platform engineering, SRE, service management, security, and compliance. Without a clear process owner and shared objectives, local pilots may improve isolated tasks while the end-to-end operational workflow remains broken. The technology may be intelligent, but the operating model still must be engineered.

Future Outlook

The future of AIOps is likely to be shaped by convergence. Observability, service management, SRE, DevOps, and AI governance are moving toward a more unified operational fabric. Instead of isolated tools for monitoring, diagnosis, automation, and review, enterprises are building closed loops that connect these stages into one managed system. This convergence supports faster learning and more coherent reliability management (Google Cloud, 2020; Microsoft, 2025c).

A second future direction is stronger evaluation discipline. Survey literature increasingly notes that many research results remain difficult to compare because datasets are synthetic, proprietary, or disconnected from production workflows. The field will mature further when evaluation focuses not only on benchmark performance but also on operational usefulness, trust, governance compliance, and business value (Notaro et al., 2021; Zhang et al., 2025).

A third future direction is selective autonomy at enterprise scale. Rather than fully autonomous operations, the more realistic model is tiered autonomy in which simple, reversible actions become highly automated while ambiguous or high-risk decisions remain human

governed. In this model, AI becomes a permanent layer of operational coordination, but not an unbounded authority.

Managerial Implications for IT Leaders

For IT leaders, the managerial implication of AIOps is that operational excellence can no longer be treated as a purely reactive support function. As digital services become central to revenue, customer experience, and regulatory exposure, incident handling, reliability engineering, and automation strategy become part of enterprise performance management. Leaders therefore need to sponsor AIOps not only as a tooling initiative but as an operating-model initiative with clear ownership, service objectives, and cross-functional alignment between engineering, service management, security, and governance.

A second implication is the need to fund foundational capabilities that may appear less visible than AI features but are more decisive in practice. High-quality observability, usable configuration and topology data, current runbooks, well-defined response policies, and disciplined post-incident review all create the conditions under which AIOps can work reliably. Organizations that skip those foundations often discover that advanced automation exposes process weaknesses rather than solving them.

A third implication is that leadership must manage the human side of operational transformation. Teams need clarity about what the system will automate, what remains under human judgment, how risk is controlled, and how success will be measured. When leaders frame AIOps as a way to remove toil, strengthen reliability, and elevate the work of engineers toward higher-value design and improvement, adoption is far more likely to be sustained.

Conclusion

AIOps represents an important shift in how organizations design and manage IT operations. Its significance lies not merely in applying machine learning or large language models to telemetry, but in redesigning the full workflow of detection, diagnosis, decision, action, and learning so that operations become faster, more consistent, and more scalable.

The analysis in this paper shows that AIOps creates the most value when four conditions are present: trustworthy observability data, strong workflow design, carefully bounded automation, and lifecycle governance. Without those conditions, AIOps can become another fragmented tool layer that increases complexity and obscures accountability. With them, it can reduce toil, improve reliability, and align IT operations more closely with business outcomes.

Therefore, AIOps should be implemented as a process-engineering capability rather than as a stand-alone AI initiative. The organizations most likely to succeed are those that treat AIOps as a managed system with explicit owners, measurable objectives, visible controls, and strong human integration. In that form, AIOps is not simply a trend in automation; it is a practical pathway to more resilient and continuously improving IT operations.

References

- Cheng, Q., Sahoo, D., Saha, A., Yang, W., Liu, C., Woo, G., Singh, M., Savarese, S., & Hoi, S. C. H. (2023). AI for IT operations (AIOps) on cloud platforms: Reviews, opportunities and challenges. arXiv. <https://doi.org/10.48550/arXiv.2304.04661>
- Diaz-de-Arcaya, J., Torre-Bastida, A. I., Zarate, G., Minon, R., & Almeida, A. (2024). A joint study of the challenges, opportunities, and roadmap of MLOps and AIOps: A systematic survey. *ACM Computing Surveys*, 56(4), Article 84. <https://doi.org/10.1145/3625289>
- Google Cloud. (2020, January 31). Identifying and tracking toil using SRE principles. <https://cloud.google.com/blog/products/management-tools/identifying-and-tracking-toil-using-sre-principles>
- Microsoft. (2024, October 23). Health modeling for workloads. Microsoft Learn. <https://learn.microsoft.com/en-us/azure/well-architected/design-guides/health-modeling>
- Microsoft. (2025a, March 6). Optimizing incident management with AIOps using the Triangle System. Microsoft Azure Blog. <https://azure.microsoft.com/en-us/blog/optimizing-incident-management-with-aiops-using-the-triangle-system/>
- Microsoft. (2025b, September 2). AIOps and machine learning in Azure Monitor. Microsoft Learn. <https://learn.microsoft.com/en-us/azure/azure-monitor/aiops/aiops-machine-learning>
- Microsoft. (2025c, December 8). Overview of Azure SRE Agent. Microsoft Learn. <https://learn.microsoft.com/en-us/azure/sre-agent/overview>
- Microsoft. (2026a, March 9). Audit agent actions in Azure SRE Agent. Microsoft Learn. <https://learn.microsoft.com/en-us/azure/sre-agent/audit-agent-actions>

Microsoft. (2026b, March 9). Track incident value in Azure SRE Agent. Microsoft Learn.

<https://learn.microsoft.com/en-us/azure/sre-agent/track-incident-value>

National Institute of Standards and Technology. (2023). Artificial intelligence risk management framework (AI RMF 1.0). U.S. Department of Commerce.

<https://doi.org/10.6028/NIST.AI.100-1>

National Institute of Standards and Technology. (2024). Artificial intelligence risk management framework: Generative artificial intelligence profile. U.S. Department of Commerce.

<https://doi.org/10.6028/NIST.AI.600-1>

Notaro, P., Cardoso, J., & Gerndt, M. (2021). A survey of AIOps methods for failure management. *ACM Transactions on Intelligent Systems and Technology*, 12(6), Article 81. <https://doi.org/10.1145/3483424>

Soldani, J., & Brogi, A. (2023). Anomaly detection and failure root cause analysis in (micro)service-based cloud applications: A survey. *ACM Computing Surveys*, 55(3), Article 59. <https://doi.org/10.1145/3501297>

Zhang, L., Jia, T., Jia, M., Wu, Y., Liu, A., Yang, Y., Wu, Z., Hu, X., Yu, P. S., & Li, Y. (2025). A survey of AIOps in the era of large language models. *ACM Computing Surveys*, 58(2). <https://doi.org/10.1145/3746635>