

SHREEVIKAS JAGADISH

Chicago, IL | (312) 358-3056 | [LinkedIn](#) | [GitHub](#) | shreevikasjagadish7@gmail.com | [Portfolio](#)

SUMMARY

AWS Certified Data Engineer and MS IT candidate with a 3.88 GPA and 2+ years of experience building cloud-native data pipelines, lakehouse models, streaming workflows, and AI data infrastructure. Hands-on with Python, SQL, PySpark, AWS, Snowflake, Databricks, Airflow, dbt, Kafka, and Great Expectations. Strong focus on reliable data ingestion, transformation, data quality, performance optimization, CI/CD, and analytics-ready datasets for business, healthcare, and AI use cases.

TECHNICAL SKILLS

Languages: Python, SQL, Spark SQL, Java, Bash/Shell Scripting, Rust

Data Engineering: PySpark, Spark Structured Streaming, Kafka, Airflow, dbt, Delta Lake, ETL/ELT, CDC, SCD Type-2, Medallion Architecture, Star Schema, Dimensional Modeling

Cloud & Warehouses: AWS S3, Glue, EMR, Lambda, Redshift, DynamoDB, Athena, SNS/SQS, CloudWatch, IAM, Snowflake, Azure Databricks

Data Quality & Reliability: Great Expectations, Data Contracts, Schema Validation, Reconciliation, Data Observability, Data Lineage, CI/CD Quality Gates

AI Data Infrastructure: LangChain, RAG, FAISS, Vector Search, OpenAI API, FastAPI, Pydantic

DevOps & Tools: Docker, Terraform, GitHub Actions, Jenkins, CI/CD, pytest, REST APIs, Structured Logging, Grafana

BI & Databases: Power BI, Tableau, PostgreSQL, MongoDB, DuckDB

EDUCATION

Master of Information Technology — Illinois Institute of Technology, Chicago, IL

August 2024 – May 2026

Bachelor of Computer Science — Visvesvaraya Technological University, Bengaluru, India

July 2019 – August 2023

WORK EXPERIENCE

AI Data Engineer Intern | NeuralSeek, Miami, FL (Remote)

July 2025 – October 2025

- Built RAG data pipelines using **LangChain**, **OpenAI embeddings**, and **FAISS** to ingest, chunk, validate, and index **10K+ real estate documents**, supporting **500+ daily queries** at sub-second latency and reducing manual knowledge lookup time by **35%**.
- Developed serverless ingestion workflows on **AWS Lambda**, **S3**, and **DynamoDB** for document processing, metadata persistence, embedding refreshes, and FAISS index updates, improving scalability for bursty document upload workloads.
- Productionized **FastAPI** services with **Pydantic schema contracts**, **structured JSON logging**, and **pytest coverage**, improving API reliability, debugging speed, and maintainability for downstream AI application teams.

Data Engineer | Whiterock Technologies, Bengaluru, India

August 2022 – April 2024

- Built **12+ cloud ELT pipelines** using **AWS S3**, **Glue**, **PySpark**, **Snowflake**, and **Airflow**, owning ingestion logic, transformation scheduling, and **Grafana/CloudWatch-based failure monitoring**, reducing reporting latency by **40%** for **200+** business stakeholders.
- Re-engineered batch ingestion into event-driven streaming workflows using **AWS SNS/SQS**, **Lambda**, **Kafka**, and **PySpark Structured Streaming**, processing **5M+ daily events** across transactions, inventory feeds, and engagement data while reducing freshness lag from hours to minutes.
- Supported **3TB+ healthcare Oracle-to-Snowflake data migration** by implementing CDC-based loads, SCD Type-2 history tracking, table-level reconciliation, checksum validation, rollback-ready cutover checks, and post-migration quality controls across **50+ tables**.
- Architected **40+ dbt** transformations using CTEs, window functions, dimensional models, and automated tests integrated with **Jenkins/GitHub Actions CI/CD quality gates**, improving metric consistency by 30% and reducing analytics deployment lead time by 25%.
- Optimized **PySpark and Snowflake workloads** on **100M+ row datasets** using partition pruning, broadcast joins, query tuning, and incremental processing, reducing compute costs by **20%** while improving pipeline runtime reliability.
- Partnered with analysts and client stakeholders to define trusted KPI datasets for sales, operations, healthcare reporting, and executive dashboards, improving dashboard consistency and reducing recurring ad-hoc reporting effort.

PROJECTS

Airbnb Analytics Lakehouse Pipeline | Snowflake, dbt, AWS S3, SQL, Python

- Built an end-to-end **Airbnb analytics lakehouse** using **AWS S3**, **Snowflake**, **dbt**, **SQL**, and **Python**, modeling raw listings, bookings, and host data through **Bronze**, **Silver**, and **Gold layers** with incremental models, dbt tests, macros, documentation, and analytics-ready tables for BI reporting.

Databricks Lakeflow Medallion Pipeline | Databricks, S3, Lakeflow Jobs, Unity Catalog, PySpark

- Developed a **Databricks Lakehouse pipeline** using **S3**, **Lakeflow Jobs**, **Unity Catalog**, **Delta/Medallion architecture**, and **PySpark notebooks** to process customer, product, and fact datasets through full and incremental loads, creating governed gold-layer tables for dashboarding and Databricks Genie-based analytics.

Stock Market Kafka Data Pipeline | Kafka, Python, AWS S3, Glue, Athena, EC2

- Built a simulated **real-time stock market streaming pipeline** using **Apache Kafka**, **Python**, **EC2**, **AWS S3**, **Glue Crawler**, **Glue Data Catalog**, and **Athena**, producing and consuming market records through Kafka topics, landing streamed data in S3, cataloging schemas with Glue, and enabling serverless SQL analytics through Athena.

CERTIFICATIONS & PROFESSIONAL DEVELOPMENT

AWS Certified Data Engineer – Associate | [Amazon Web Services](#)

Databricks Certified Associate Developer for Apache Spark | In Progress (Expected May 2026)