

SHREEVIKAS JAGADISH

Chicago, IL | (312) 358-3056 | [linkedin.com/in/shreevikasbj](https://www.linkedin.com/in/shreevikasbj) | github.com/Shreevikas-BJ | shreevikas-portfolio.vercel.app | shreevikasjagadish7@gmail.com

SUMMARY

AWS Certified Data Engineer with 2+ years of experience building batch, streaming, and AI-ready data pipelines across AWS, Snowflake, Databricks, Airflow, dbt, Kafka, and PySpark. Skilled in CDC, data quality, Spark optimization, RAG workflows, BI/ML-ready datasets, with research experience in Physics-Informed Neural Networks using NVIDIA PhysicsNeMo.

TECHNICAL SKILLS

Languages: Python, SQL, Spark SQL, Bash

Data Engineering: PySpark, Kafka, Airflow, dbt Core, Fivetran, ETL/ELT, CDC, SCD Type-2

Cloud & Warehouses: AWS S3, Glue, EMR, Lambda, Snowflake, Redshift, Azure Databricks, SageMaker

Lakehouse & Modeling: Delta Lake, Medallion Architecture, Unity Catalog, Dimensional Modeling

Data Quality & Reliability: Great Expectations, dbt Tests, Schema Validation, Reconciliation, Data Lineage, SLA Monitoring

AI/ML: RAG Pipelines, Embeddings, Vector Search, FAISS, Scikit-Learn, XGBoost, LightGBM, MLflow, LangChain, OpenAI API, Codex

DevOps & BI: Docker, GitHub Actions, pytest, Grafana, Power BI, Tableau

EDUCATION

Master of Information Technology - Illinois Institute of Technology, Chicago, IL

August 2024 - May 2026

Bachelor of Computer Science - Visvesvaraya Technological University, Bengaluru, India

July 2019 - August 2023

WORK EXPERIENCE

AI Data Engineer Intern | NeuralSeek, Miami, FL (Remote)

July 2025 - November 2025

- Led a 4-member intern team to design a LangChain-based RAG architecture build ingestion, chunking, embedding, and FAISS indexing pipelines for 10K+ real estate documents, supporting 500+ daily RAG queries.
- Improved document upload reliability by 30% by building Lambda, S3, and DynamoDB workflows for metadata tracking, embedding refreshes, vector index updates, and MCP-based system-error notifications.
- Reduced UI/API iteration time by 25% using Codex-assisted development to accelerate FastAPI updates, Pydantic schema validation, structured logging, automated test runs, pytest coverage, and error handling.

Data Engineer | Whiterock Technologies, Bengaluru, India

August 2022 - April 2024

- Reduced reporting latency by 40% for 200+ users by owning 12+ AWS-to-Snowflake ELT pipelines with Fivetran ingestion, Airflow scheduling, dbt tests, 40+ dbt Core models, and failure monitoring.
- Improved active dashboarding from hourly batch updates to near-real-time views by implementing Kafka-to-Spark streaming workflows across 8 Kafka topics, processing ~5M daily transaction and engagement events.
- Operationalized 2 ML pipelines using XGBoost and LightGBM with PySpark feature engineering, SageMaker batch inference, Great Expectations quality gates, Grafana monitoring, and 98.8% uptime via Airflow and CloudWatch.
- Led a 3TB+ Oracle-to-Snowflake migration across 50+ healthcare tables by rewriting procedural logic, validating CDC/SCD Type-2 loads, and preserving role-based access controls.

Data Engineer Intern | Whiterock Technologies, Bengaluru, India

February 2022 - July 2022

- Reduced manual reporting effort by automating extraction and validation tasks with Python, SQL, and Excel while supporting Airflow retries and failure alerts.
- Improved reporting accuracy and query performance by writing optimized SQL with CTEs, window functions, joins, and aggregations for analyst-facing datasets.
- Supported healthcare migration readiness through schema validation, row-count reconciliation, checksum checks, mismatch investigation, and post-load documentation.

PROJECTS

Databricks Lakeflow Medallion Pipeline | github.com/Shreevikas-BJ/databricks-lakeflow-medallion-pipeline

- Built a Databricks Medallion pipeline using S3, Delta tables, PySpark, Lakeflow Jobs, and Unity Catalog.
- Processed customer, product, and transaction datasets through Bronze, Silver, and Gold lakehouse layers.
- Improved analytics access by creating cataloged Gold-layer tables for dashboards, business queries, and Databricks Genie.

ML Course Document RAG | github.com/Shreevikas-BJ/ml-course-document-rag

- Built a document-grounded RAG system for ML course PDFs using text extraction, chunking, embeddings, and FAISS vector search.
- Improved answer reliability with similarity threshold gating, citation-based responses, and refusal handling.
- Added a Multi-LLM Council review layer to cross-check retrieved context quality before final response generation.

CERTIFICATIONS

AWS Certified Data Engineer - Associate | verify: [credly.com/badges/017bc7a0-a378-4cfa-abb0-bc968c20d7da](https://www.credly.com/badges/017bc7a0-a378-4cfa-abb0-bc968c20d7da)

Google Data Analytics Specialization Certificate | verify: coursera.org/account/accomplishments/specialization/TIT1TAQNFGPT