

SHREEVIKAS JAGADISH

Chicago, IL | (312) 358-3056 | [linkedin.com/in/shreevikasbj](https://www.linkedin.com/in/shreevikasbj) | github.com/Shreevikas-BJ | shreevikas-portfolio.vercel.app | shreevikasjagadish7@gmail.com

SUMMARY

AWS Certified Data Engineer with 2+ years of experience building batch, streaming, and AI-ready data pipelines across AWS, Snowflake, Databricks, Airflow, dbt Core, Kafka, and PySpark. MS in Information Technology graduate from Illinois Institute of Technology, skilled in data ingestion, CDC, dimensional modeling, data quality, Spark optimization, and analytics-ready datasets for BI and ML use cases.

TECHNICAL SKILLS

Languages: Python, SQL, Spark SQL, Bash

Data Engineering: PySpark, Kafka, Airflow, dbt Core, Fivetran, ETL/ELT, CDC, SCD Type-2

Cloud & Warehouses: AWS S3, Glue, EMR, Lambda, Snowflake, Redshift, Azure Databricks

Lakehouse & Modeling: Delta Lake, Apache Iceberg, Medallion Architecture, Unity Catalog, Dimensional Modeling

Data Quality & Reliability: Great Expectations, dbt Tests, Schema Validation, Reconciliation, Data Lineage, SLA Monitoring

AI/ML Infrastructure: RAG Pipelines, Embeddings, Vector Search, FAISS, LangChain, OpenAI API, FastAPI

DevOps & BI: Docker, GitHub Actions, Jenkins, pytest, Grafana, Power BI, Tableau

EDUCATION

Master of Information Technology - Illinois Institute of Technology, Chicago, IL

August 2024 - May 2026

Bachelor of Computer Science - Visvesvaraya Technological University, Bengaluru, India

July 2019 - August 2023

WORK EXPERIENCE

AI Data Engineer Intern | NeuralSeek, Miami, FL (Remote)

July 2025 - November 2025

- Led a 4-member intern team at NeuralSeek, an AI knowledge automation platform, and supported 500+ daily RAG queries by building ingestion, chunking, embedding, and FAISS indexing pipelines for 10K+ real estate documents.
- Improved upload reliability by 30% by using Lambda, S3, and DynamoDB for metadata storage, embedding refreshes, vector index updates, and RAG pipeline tuning.
- Reduced UI/API iteration time by 25% by using Codex-assisted development to build the application interface and improve FastAPI services, Pydantic schemas, logging, tests, and error handling.
- Improved debugging efficiency by 35% by automating test runs, MCP-based system-error notifications, and Excel-based user-query feedback logs for recurring issue analysis.

Data Engineer | Whiterock Technologies, Bengaluru, India

August 2022 - April 2024

- Reduced reporting latency by 40% for 200+ users by owning 12+ AWS-to-Snowflake ELT pipelines with Fivetran ingestion, Airflow scheduling, dbt tests, Apache Iceberg tables, and failure monitoring.
- Improved active dashboarding from hourly batch updates to near-real-time views by implementing Kafka-to-Spark streaming workflows across 8 Kafka topics, processing ~5M daily transaction and engagement events at ~3K events/minute.
- Led a 3TB+ Oracle-to-Snowflake migration across 50+ healthcare tables by rewriting procedural logic, validating CDC/SCD Type-2 loads, and replicating Oracle user roles.
- Improved metric consistency by 30%, reduced compute overhead by 20%, and maintained 99.8% pipeline SLA by developing 40+ dbt Core models, optimizing Spark/Snowflake workloads, and owning monitoring for production reporting workflows.

Data Engineer Intern | Whiterock Technologies, Bengaluru, India

January 2022 - July 2022

- Reduced manual reporting effort and improved pipeline reliability by automating extraction and validation tasks with Python, SQL, and Excel while supporting Airflow retries, failure alerts, and CloudWatch/Grafana monitoring.
- Improved query performance and reporting accuracy by writing optimized SQL using CTEs, window functions, joins, and aggregations for analyst-facing datasets.
- Strengthened healthcare migration readiness by supporting schema validation, row-count reconciliation, checksum checks, mismatch investigation, and post-load validation documentation.

PROJECTS

Databricks Lakeflow Medallion Pipeline | github.com/Shreevikas-BJ/databricks-lakeflow-medallion-pipeline

- Built a Databricks Medallion pipeline using S3, Delta tables, PySpark, Lakeflow Jobs, and Unity Catalog.
- Processed customer, product, and transaction datasets through Bronze, Silver, and Gold lakehouse layers.
- Improved analytics access by creating cataloged Gold-layer tables for dashboards, business queries, and Databricks Genie.

ML Course Document RAG | github.com/Shreevikas-BJ/ml-course-document-rag

- Built a document-grounded RAG system for ML course PDFs using text extraction, chunking, embeddings, and FAISS vector search.
- Improved answer reliability with similarity threshold gating, citation-based responses, and refusal handling.
- Added an LLM Council review layer to cross-check retrieved context quality before final response generation.

CERTIFICATIONS

AWS Certified Data Engineer - Associate | verify: [credly.com/badges/017bc7a0-a378-4cfa-abb0-bc968c20d7da](https://www.credly.com/badges/017bc7a0-a378-4cfa-abb0-bc968c20d7da)

Google Data Analytics Specialization Certificate | verify: coursera.org/account/accomplishments/specialization/TIT1TAQNFGPT