

SHREEVIKAS JAGADISH

Data Science | Machine Learning | Predictive Modeling | Statistical Learning | Decision Intelligence
United States | shreevikas@workapplyjob.com | +1 (210) 201-6775 | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

Data Scientist with 4 years of experience translating complex business challenges into scalable analytical solutions using machine learning, statistical modeling, and cloud based data platforms. Experienced in developing production ready predictive systems, optimizing large scale data workflows, and delivering measurable operational improvements through robust model deployment, experimentation, and cross functional collaboration across enterprise technology environments.

PROFESSIONAL WORK EXPERIENCE

Cloudera, United States

January 2026 – Present

Data Scientist / AI-ML

- Designed regional demand forecasting models using Python, SQL, PySpark, and XGBoost to capture product level seasonality and demand patterns, improving inventory planning accuracy by 24% across distributed supply operations.
- A reusable experimentation framework built with Apache Spark and MLflow shortened model iteration cycles by 40% while standardizing feature processing, training workflows, experiment comparison, version tracking, and production releases.
- Integrated predictive inference services and LLM powered workflows through FastAPI and Docker, cutting processing latency by 38% while enabling business applications to deliver timely operational recommendations and contextual insights.
- To strengthen production reliability, established statistical validation, retrieval evaluation, drift monitoring, and A/B testing practices that increased ML and RAG performance consistency by 18% across evolving data distributions.
- Supply chain planning and enterprise knowledge discovery requirements were translated into deployable AI solutions through close collaboration with analytics, product, and operational stakeholders throughout development and implementation cycles.
- Executive teams gained clearer visibility into forecast performance, model predictions, operational exceptions, and planning trends through Power BI dashboards that eliminated 60% of recurring manual reporting effort.

Bosch, India

March 2021 – July 2024

Data Scientist

- Engineered manufacturing quality prediction models with Python and Scikit Learn to identify complex production patterns, raising defect detection accuracy by 22% across high volume manufacturing lines and operational environments.
- Manufacturing datasets and unstructured maintenance records were consolidated through automated ETL and NLP preprocessing workflows, cutting preparation time by 45% while improving analytical consistency for downstream modeling activities.
- Improved prediction precision by 17% across quality and predictive maintenance use cases through systematic feature selection, hyperparameter tuning, cross validation, and optimization of classification and deep learning models.
- Statistical analysis and anomaly detection dashboards provided engineering teams with earlier visibility into process instability, recurring production patterns, and equipment behavior requiring further investigation or corrective action.
- Repetitive analytics and operational reporting processes were automated using SQL, Pandas, and scheduled data pipelines, saving 55% of manual processing effort while delivering more consistent insights to manufacturing leadership.

PROJECTS

Intelligent Inventory Demand Forecasting Platform | Python, PySpark, XGBoost, MLflow, FastAPI, AWS

- Built a demand forecasting platform using PySpark and XGBoost to analyze multi regional inventory patterns, improving forecast accuracy by 21% and supporting product level replenishment planning.
- Automated model training, MLflow experiment tracking, performance validation, and FastAPI deployment through AWS pipelines, reducing model release cycles by 35% and accelerating forecast delivery.

Manufacturing Process Quality Intelligence System | Python, Scikit Learn, SQL, Pandas, PostgreSQL, Tableau

- Developed ML and deep learning models to analyze manufacturing variables and detect abnormal process behavior, improving early defect identification by 19% across production quality workflows.
- Integrated automated data preparation, anomaly detection, statistical validation, and Tableau reporting, reducing manual analysis time by 42% while improving visibility into recurring production issues.

Enterprise Knowledge Search Platform | Python, LangChain, Hugging Face, Vector Search, FastAPI, PostgreSQL

- Engineered a RAG based knowledge platform using transformer embeddings, semantic chunking, and vector retrieval, improving relevant document discovery by 31% across internal technical knowledge repositories.
- Implemented prompt orchestration, retrieval evaluation, grounded response generation, and FastAPI inference services, reducing information search time by 46% while improving contextual access to enterprise documentation.

EDUCATION

Master of Science in Information Technology & Management | *Illinois Institute of Technology* | USA

May 2026

Bachelor in Computer Science | *Visvesvaraya Technological University* | India

August 2023

SKILLS

Programming: Python, SQL, PySpark, R

Machine Learning: Supervised & Unsupervised Learning, Feature Engineering, Model Optimization, Classification, Regression, Clustering, Time Series Forecasting, Recommendation Systems, Anomaly Detection

Generative AI: Large Language Models (LLMs), RAG, Hugging Face, Transformers, Prompt Engineering, Embedding Models, LangChain

Deep Learning: TensorFlow, PyTorch, Scikit Learn, XGBoost, LightGBM

Data Engineering: Apache Spark, Pandas, NumPy, ETL Pipelines, Data Processing, Data Quality

MLOps & Deployment: MLflow, Docker, Kubernetes, FastAPI, Model Serving, Model Monitoring, CI/CD

Cloud & Data Platforms: AWS (S3, SageMaker, EC2, Lambda), Azure, Databricks, Snowflake, BigQuery

Visualization & Analytics: Power BI, Tableau, Matplotlib, Plotly, Statistical Analysis, A/B Testing

Databases & Tools: PostgreSQL, SQL Server, MySQL, MongoDB, Git, Linux, Jira, Confluence, REST APIs, Agile, Scrum